

Deploy K8s Cluster with DeepOps and Kubespray

Full Documentation:

- Read full Kubernetes on DeepOps documentation [here](#).
- Read full Kubespray Getting Started documentation [here](#).


Minimum Specifications:

3 Master Nodes (VM's) soft limits

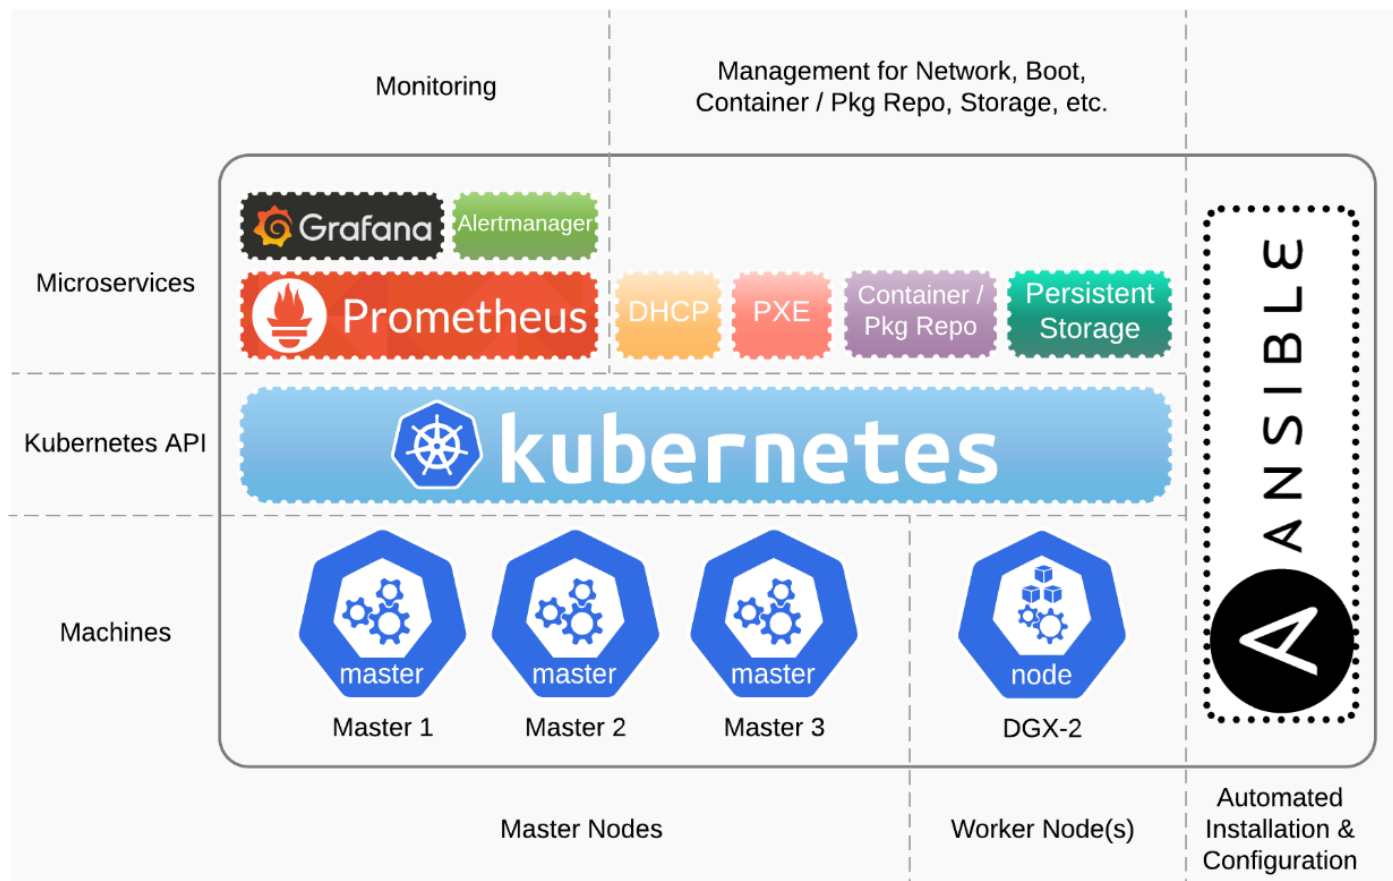
- CPU: 4 CPU cores
- Memory: 32G
- Disk: 100G
- Network: 2 NIC's
- OS: Ubuntu 18.04 LTS

Table of Contents:


Full Documentation	1
Minimum Specifications	1
DeepOps Architecture Overview.....	2
Possible System Architecture (Cisco)	3
Possible Network Architecture (Cisco)	3
Create Cluster:	4
Worker Nodes:	4
Main Master Node:	4
Troubleshooting	4
Install Cluster Add-ons:.....	5
Main Master Node:	5
Scale Cluster.....	7
Add node.....	7
Remove node	7
Delete Cluster.....	7
Note	7
Appendix A: Compatibility	8
Appendix B: Example Inventory File.....	9
Appendix C: Example Hosts File	11



DeepOps Architecture Overview:

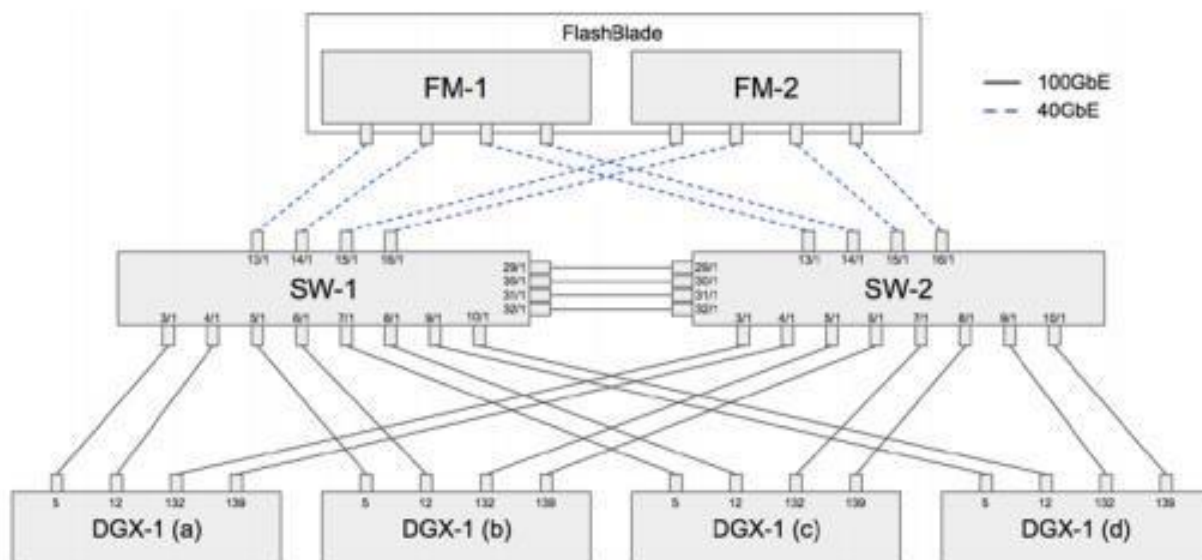


Possible System Architecture (Cisco):



Source: AIRI Reference Architecture with Cisco Nexus 9300 Ethernet Switch

Possible Network Architecture (Cisco):



Source: AIRI Reference Architecture with Cisco Nexus 9300 Ethernet Switch

Create Cluster:

Worker Nodes:

- Add SSH keys from master node(s) to worker node(s) with the command:

```
ssh-copy-id -i ~/.ssh/id_rsa.pub user@host
```

Main Master Node:

- Clone DeepOps [repo](#) and enter directory.
- To setup provisioning machine, run: `./scripts/setup.sh`.
- Run `cp -r config.example config` and edit `config/inventory` to add nodes.
- Format like: `node01 ansible_host=127.0.0.1 ip=127.0.0.1`
 - See Appendix B for an example.
- Create cluster via Ansible/Kubespray with the command: `ansible-playbook -i config/inventory playbooks/k8s-cluster.yml -k -K`

Troubleshooting:


- During installation:
 - Error: Failed to lock apt for exclusive operation:
 - Continue anyways, it should work regardless.
 - Error: Could not get lock /var/lib/dpkg/lock-frontent - open (11: Resource temporarily unavailable)
 - See [this](#) AskUbuntu page if the error causes problems with the installation. Even though it says nodes failed, the cluster should come up without a problem.
- Post-installation:
 - coredns won't deploy: "0/X nodes are available: X node(s) didn't match pod affinity/anti-affinity, X node(s) didn't satisfy existing pods anti-affinity rules"
 - The deployment is likely trying to create X+1 pods, where X is the number of nodes available. Scale the deployment down to the number of available nodes, and the error should go away.
 - tiller-deploy won't deploy: "0/X nodes are available: X node(s) didn't match node selector, X node(s) had taints that the pod didn't tolerate"
 - Remove the taint on one of your master nodes with `kubectl taint node <node-name> node-role.kubernetes.io/master:NoSchedule-`




Install Cluster Add-ons:

Main Master Node:

- Install GPU drivers with `ansible-playbook -i config/inventory playbooks/k8s-gpu-plugin.yml -K -k`
 - **Troubleshooting:** If you cannot schedule pods with GPU resources, check the log of the `nvidia-device-plugin-daemonset` pod on the GPU node to find the error.
 - If the log says “Failed to initialize NVML: could not load NVML library”:
 - Check to make sure `/etc/docker/daemon.json` on your GPU node is configured to use the `nvidia-container-runtime`.
 - If the docker runtime is configured correctly and you still cannot schedule GPU resources, try the following commands on your GPU node:
 - `systemctl daemon-reload`
 - `systemctl restart docker`
 - If you still cannot schedule pods with GPU resources:
 - Try running `ansible-playbook -i config/inventory playbooks/nvidia-driver.yml -K -k` to manually install the Nvidia GPU drivers.
 - If all else fails, remove the GPU node or tear down the entire cluster and rebuild it.
 - Install K8s dashboard with `./scripts/k8s_deploy_dashboard_user.sh`
 - Script will:
 - Install dashboard and create any resources required (including deployments, services, service accounts, etc.)
 - Automatically expose an endpoint for the service (port changes per script run).
 - Output a URL to open in web browser.
 - Output a token to access K8s dashboard login page.
 - Install monitoring with `./scripts/k8s_deploy_monitoring.sh`
 - Script will:
 - Install dashboard and create any resources required (including deployments, services, service accounts, etc.)
 - Automatically expose an endpoint for the service (ports are set).
 - Output URLs to open in web browser to access dashboards.
 - Services can be reached from:
 - Grafana: <http://mgmt:30200>



- Prometheus: <http://mgmt:30500>
- Alertmanager: <http://mgmt:30400>
- Install persistent storage with `./scripts/k8s_deploy_rook.sh`
 - Script will:
 - Install Rook using Ceph as a backend.
 - Deploy Ceph file and block storage.
 - Deploy a Ceph cluster manager and toolbox container.
 - Deploy the Ceph dashboard.
 - Use `./services/rook-cluster.yml` as the deployment configuration – make any desired changes here.
 - **Troubleshooting:** If you've previously installed rook-ceph and your rook-ceph pods are not deploying, try deleting everything in the default data directory with `sudo rm -rf /var/lib/rook` on all nodes.
 - To remove the Rook deployment, run `./scripts/rmrook.sh`
- Install Kubeflow with `./scripts/k8s_deploy_kubeflow.sh`
 - Full Kubeflow documentation can be read [here](#).
 - Script will:
 - Install dashboard and create any resources required (including deployments, services, service accounts, etc.)
 - Automatically expose an endpoint for the service (port changes per script run).
 - Output a URL to open in web browser to access Kubeflow dashboard.



Scale Cluster:

- Recreate inventory file using the inventory builder tool as seen [here](#) under Quick Start (see Appendix C for an example).
- Make sure to only have Kubernetes resources specified here such as `[all]`, `[kube-master]`, `[kubenode]`, `[etcd]`, and `[k8s-cluster:children]`.
- Additionally, make sure these resources match the previous inventory file.
- In `kubespray/` directory:

Add node:

- Add new node to `inventory/mycluster/hosts.yml` (as well as `../config/inventory`).
- Run `ansible-playbook -i inventory/mycluster/hosts.yml --become --become-user=root scale.yml -K -k`

Remove node:


- Run `ansible-playbook -i inventory/mycluster/hosts.yml --become --become-user=root remove-node.yml --extra-vars "node=node-to-remove" -K -k`
- Remove node from inventory files.

Delete Cluster:

- Recreate inventory file using the inventory builder tool like before.
- Additionally, make sure these resources match the previous inventory file.
- Script will:
 - Fully tear down cluster.
 - Uninstall kubectrl, kubeadm, kubelet from nodes.
- In `kubespray/` directory:
 - Run `ansible-playbook -i inventory/mycluster/hosts.yml --become --become-user=root reset.yml -K -k`

Note:


- If during Cluster deletion or Node removal, the Ansible script hangs on “TASK [reset : reset | unmount kubelet dirs”, try restarting docker service on all nodes via `systemctl restart docker.service` .



Appendix A: Compatibility

DeepOps Version: 19.07

Software	Version
Ansible	2.8.11
Kubespray	v2.10.4
Kubernetes	v1.14.3
Docker	18.09.6
Rook	v1.0.2
Ceph	v13 (v13.2.6-20190604)
Slurm	19.05



Appendix B: Example Inventory File

deepops/config/inventory

```
# Server Inventory File
#
# Uncomment and change the IP addresses in this file to match your environment
# Define per-group or per-host configuration in group_vars/*.yaml

#####
# ALL NODES
#####
[all]
mgmt-node01    ansible_host=127.0.0.1 ip=127.0.0.1
worker-node01  ansible_host=127.0.0.2 ip=127.0.0.2
worker-node02  ansible_host=127.0.0.3 ip=127.0.0.3
#####
# KUBERNETES
#####
[kube-master]
mgmt-node01    ansible_host=127.0.0.1 ip=127.0.0.1

[etcd]
mgmt-node01    ansible_host=127.0.0.1 ip=127.0.0.1


[kube-node]
worker-node01  ansible_host=127.0.0.2 ip=127.0.0.2
worker-node02  ansible_host=127.0.0.3 ip=127.0.0.3

[k8s-cluster:children]
kube-master
kube-node

#####
# SLURM
#####
[slurm-master]

[slurm-node]

[slurm-cluster:children]
```



```
#####
# NFS
#####
[nfs-server]

[nfs-clients]

#[nfs-clients:children]

#####
# OFFLINE CACHE BUILDER
#####
[cache_builder]
localhost  ansible_connection=local

#####
# SSH connection configuration
#####
[all:vars]
#ansible_user=ubuntu
# Configure SSH bastion/jumpbox for the cluster
#ansible_ssh_common_args='-o ProxyCommand="ssh -W %h:%p -q ubuntu@10.0.0.1"'
```



Appendix C: Example Hosts File

deepops/kubespray/inventory/mycluster/hosts.yml

```
[all]
mgmt-node01      ansible_host=127.0.0.1 ip=127.0.0.1
worker-node01    ansible_host=127.0.0.2 ip=127.0.0.2
worker-node02    ansible_host=127.0.0.3 ip=127.0.0.3
```

```
[kube-master]
mgmt-node01
```

```
[kube-node]
worker-node01
worker-node02
```

```
[etcd]
mgmt-node01
```

```
[k8s-cluster:children]
kube-node
kube-master
```

